

Mastering the game of Go with deep neural networks and tree search

Summary: This work has introduced a new algorithm that combines Monte Carlo simulation with value and policy networks to be utilized in GO game and that defeated for the first time a human champion!

Introduction:

- **GO** is the most challenging classic game because:
 - It has an enormous search space.
 - Difficulty of evaluating board positions and moves
- Given infinite resources, the optimal move of a deterministic game is completely determined by recursively searching a tree with possible b^d moves, where b is the number of legal moves and d is the game length.
- For games with large playing space like **GO**, where b^d soon becomes intractable, certain strategies were introduced to limit the search space and yet produce a reasonable approximation. These include:
 1. Reducing the depth by position evaluation with an approximate scoring functions for deep subtrees.
 2. Sampling actions over possible moves probability distribution, and maybe averaging the outcome.
- Because of the complexity of the **GO** game, only the second solution may provide an acceptable performance, if the sampling strategy was enhanced. In that regard, **Monte Carlo Tree Search (MCTS)** is one of the widely utilized techniques to predict human expert moves.

Methodology:

- Utilizing convolutional networks by passing an image board of size 19x19 and allowing the deep network to abstract various information. They constructed two somehow similar networks:
 - "value networks" to evaluate board positions. Outputs probabilities
 - "Policy networks" to select moves. It outputs a single prediction
- **Training Methodology**
 - supervised learning from human expert games:
 - Series of convolutional layers and nonlinear rectifiers with final softmax layer that outputs the probabilities over all legal moves
 - Stochastic gradient decent was utilized to maximize the likelihood of human move given a selected state.
 - They achieved up to 57% accuracy using all input features.
 - reinforcement learning from games of self-play.
 - To prevent overfitting, the game was played between current selected policy network and a randomly selected one by utilizing a certain rewarding function.
 - Weights are updated using stochastic gradient decent to maximize the expected outcome.
 - For certain games, they achieved around 85% win.
 - They also plotted the MSE value between the predicted value and the actual game.
- **Searching**
 - Combines policy and value networks in an MCTS algorithm and select actions by lookahead search.
 - The tree is traversed by simulation, starting from the root state.
 - An action is selected in accordance with some prior probability and a visit count for all traversed edges is kept updated.
 - Utilizes multi-threaded search that executes simulations on CPUs and computes policy and value networks in parallel on GPUs.

Results

- They tested their model against various commercial as well as open source programs that already utilizes high-performance MCTS algorithms with a time limit of **5 sec** to compute the next move.
- They reported a winning scenario in 99.8% against other go programs.
- They tried different challenges and their program was winning most of the time.
- They reported that their distributed version is much stronger.
- And for the first time in this game's history, they reported defeating the human European Go champion by 5 games to 0.
- Their algorithm, chooses next moves intelligently and thus requires searching less space

Achievements

- By combining tree search with policy and value networks, the introduced a frame work where human -level performance may be achieved in intractable AI domains. The method paves the way to advances in many other domains like
 - General game playing
 - Classical planning
 - Partially observed planning
 - Scheduling
 - Constraint Satisfaction

Remarks

- **I liked to see how this algorithm would compete with other commercial products if the time per move was extended to like 30s rather than 5. 😊**